

Concurrency Control

lock-based protocols

- To ensure serializability is to require that data items be accessed in a mutually exclusive manner, i.e. while one transaction is accessing a data item, no other transaction can modify that data item.
- The most common method used to implement this requirement is to allow a transaction to access data item only if it is currently holding a lock on that item.

Locks-
→ There are various modes in which data item can be locked.

① Shared lock- If a transaction T_i has obtained a shared-mode lock (denoted by S), on data item B_j , then T_i can read, but cannot write.

② Exclusive lock- If a transaction T_i has obtained an exclusive-mode lock (denoted by X) on data item B_j , then T_i can perform both read and write.

- A transaction requests a shared lock on data item A by executing the lock-S(A).
- A transaction request a exclusive lock through the lock-X(A).
- A transaction ~~can unlock~~ a data item A by unlock(A). instruction.

T₁ :- lock-X(B);
 read(B);
 B := B - 50;
 write(B);
 unlock(B);
 lock-X(A);
 read(A);
 A := A + 50;
 write(A);
 unlock(A).

T₂ :- lock-S(A);
 read(A);
 unlock(A);
 lock-S(B);
 read(B);
 unlock(B);
 display(A+B).

transaction T₂.

Transaction T₁ must access a data item, transaction T₁ must

- To access a data item, first lock that item. If the data item is already locked by another transaction in a incompatible mode, the concurrency control manager will not grant a lock until all incompatible locks held by another transaction have been released.

Concurrently control Manager

T_1

lock - X (B)

read (B)
 $B := B + 50$
write (B)
unlock (B)

grant - X (B, T_1)

unlock (A)

lock - S (A)

read (A)
unlock (A)
unlock (B)

grant - S (A, T_2)

read (B)
unlock (B)
display (A+B)

lock - X (A)

read (A)
 $A := A + 50$
write (A)
unlock (A)

grant - X (A, T_2)

(a)

The above table explains how the concurrency control manager is assigned to locks and released the locks b/w two transactions T_1 & T_2 .

Granting of locks:-

- suppose a transaction T_2 has a shared-mode lock on data item,
- Another transaction T_1 requests an exclusive mode lock on the same data item.
- T_1 has to wait for T_2 to release the shared-mode lock.
- Meanwhile, a transaction T_3 may request a shared-mode lock on the same data item.
- The lock request is compatible with the lock granted to T_2 , so T_3 may be granted shared-mode lock.
- At this point T_2 may release the lock, but still T_1 has to wait for T_3 to finish.
- Suppose T_3 new transaction requests a shared mode lock on the same data item, and it is granted the lock before T_3 is released.

→ In this process T_1 never gets the chance of exclusive-mode lock on the data item.

→ This situation is called starvation.

→ To avoid starvation of transactions by granting locks in the following manner.

⇒ When a transaction T_i requests a lock on data item α in a particular mode M , the concurrency control Manager grants lock provided that

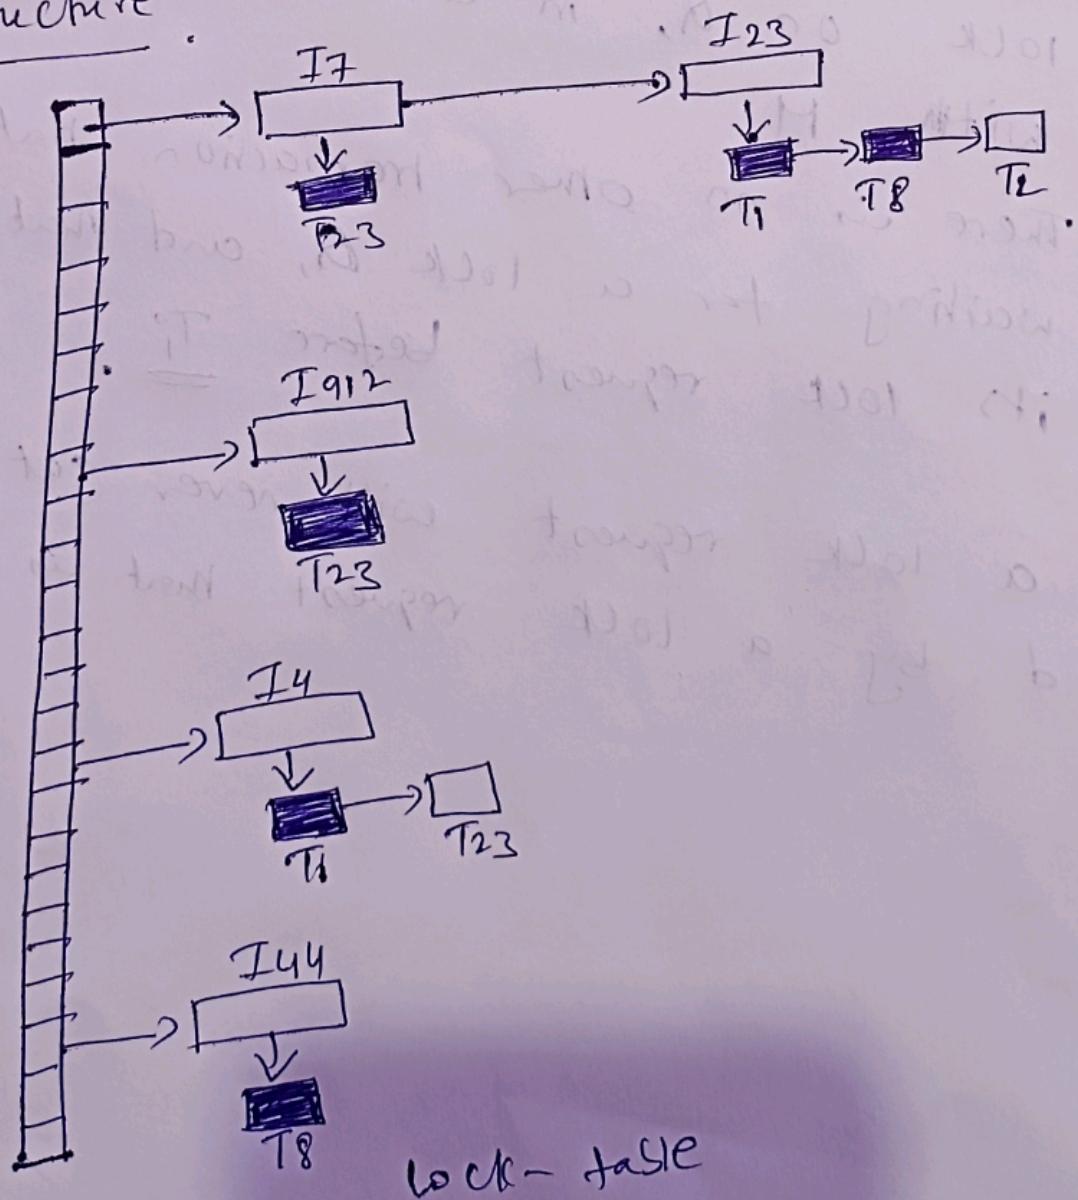
① There is no other transaction holding a lock on α , in a mode that conflicts with M .

② There is no other transaction waiting for a lock α , and that made its lock request before T_i .

Thus, a lock request will never get blocked by a lock request that is made later.

Implementation of locking:

- A lock manager can be implemented as a process that receives messages from transactions and sends messages in reply.
- The lock manager replies to lock-request messages with lock-grant messages, or with messages requesting rollback of the transaction (in case of deadlock).
- The lock manager uses the following data structure.



(7)

- for each data item that it currently locked, it maintains a linked list of records.
- one for each request, in the order in which the requests arrived.
- It uses a hash table indexed on the name of data item.
- To find the linked list if any for a data item.
- This table is called the lock-table.
- Each record of the linked list for a data item notes which transaction made the request, and what lock mode is requested.
- The record also notes if the request has currently been granted.

In the above lock-table.

- Table contains locks for 5 different data items.
- I₄, I₇, I₂₃, I₄₄ and I₉₁₂.
- lock-table uses overflow chaining.
- There is a linked-list of data items for each entry in the lock-table.
- There is also a list of transactions that have been granted locks, waiting for locks.

→ Granted locks are filled in black rectangles



→ waiting request are the empty rectangles. □.

→ T_{23} has been granted locks on I_{912} and I_7 . and it is waiting for I_4 .

lock-Manager processes requests this way

→ When a lock request message arrives, it adds a record to the end of the linked list for the data item, if the linked list is present.

→ Otherwise it creates a new linked list containing only the record for the request.

→ It always grants the first lock request on a data item.

→ If the transaction requests a lock on a item on which a lock has already been granted, the lock manager grants the request only if it is compatible with all earlier requests.

→ If the earlier requests have been granted already.

→ Otherwise the request has to wait until lock is released.

Validation based protocols-

→ Each transaction T_i executes in two or more different phases in its lifetime, depending on whether it is read-only or an update transaction.

→ The phases are

① Read Phase- During this phase, the system executes transaction T_i .

→ It reads the value of the various data items and stores them in variables local to T_i .

→ It performs temporary local updates of the database.

② Validation Phase- Transaction T_i performs a validation test to determine whether

it can copy to the database that hold the temporary local variables.

results of write operation. without causing violation of serializability.

- ③ Write phase:- If transaction T_i succeeds in validation step 2 done,
 → The system applies the actual updates to the database.
- The system rolls back T_i .
- Each transaction must go through three phases,

To perform the validation test, transaction T_i has 3 different timestamps.

① Start (T_i), the time when T_i started its execution.

② Validation (T_i), the time when T_i finished its read phase and started its validation.

③ Finish (T_i), the time when T_i finished its write phase.

The validation test for T_j requires that, for all transactions T_i with $T_i \neq T_j$, $TS(T_i) < TS(T_j)$ must hold one of the 2 conditions

① $\text{Finish}(T_i) < \text{start}(T_j)$.

→ since T_i completes its execution before T_j started, the serializability order is indeed maintained.

② The set of data items written by T_i does not intersect with the set of data items read by T_j , and T_i completes its write phase before T_j starts its validation phase.

$\boxed{\text{start}(T_j) < \text{Finish}(T_i) < \text{validation}(T_j)}$