

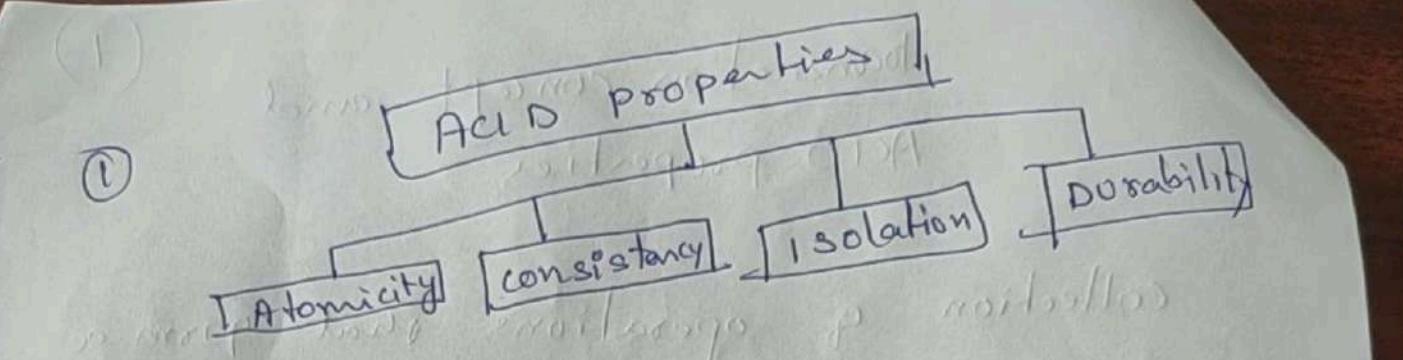
Transaction concept and ACID properties

collection of operations that form a single logical unit of work are called transaction.

A transaction is a unit of program execution that accesses and possibly update various data items.

The transaction consist of all operations executed between the begin transaction and end transaction.

To ensure integrity of the data, the database system maintain the following properties of the transaction. These properties are called ACID properties.



① Atomicity! This means that either the entire transaction takes place at once or doesn't happen at all. Transactions do not occur partially. Each transaction is considered as one unit and either goes to completion or is not executed at all. It involves the following two operations.

- Abort! - If a transaction aborts, changes made to the database are not visible.
- Commit! - If a transaction commits, changes made are visible.

Atomicity is also known as
the "All or nothing rule"

Consider the following transaction T
consisting of T₁ and T₂: Transfer
of 100 from account X to account Y.

Before X: 500		Y: 200	
Transaction T			
T ₁			T ₂
Read(X)			Read(Y)
X := X - 100			X := Y + 100
Write(X)			Write(Y)
			Y: 300
After X: 400			

If the transaction fails after
completion of T₁ but before completion
(or) of T₂, say after write(X) and
before write(Y)) then the
amount has been deducted
from X but not added to Y.

This results in an inconsistent
database state.

Therefore, the transaction must be executed in its entirety in order to ensure the correctness of the database state.

→ consistency, - This means that Integrity constraints must be maintained so that the database is consistent before and after the transaction. It refers to the correctness of a database.

Referring the example above,

The total amount before and after the transaction must be maintained.

Total before T occurs - 500 +
200 = 700

Total after T occurs - 400 +
300 = 700

Therefore, the database is consistent.

3

in consistency, occurs in case T_1 completes but T_2 fails. As a result, T is incomplete.

Isolation :- Even though multiple transactions can occur simultaneously, the system guarantees that for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started or T_j started execution after T_i finished. Thus, each transaction is unaware of other transactions executing concurrently in the system.

This property ensures that the execution of transactions concurrently will result in a state that is

equivalent to a state achieved where
they were executed serially in some
order.

Let $x = 500$, $y = 500$

consider two transaction T_1 and
 T_2 .

T_1	T_2
Read(x)	Read(x)
$x' = x * 100$	Read(y)
write(x)	$z = x + y$
Read(y)	write(z)
$y' = y - 50$	
write(y)	

Suppose T_1 has been executed till
Read(x) & and then T_2
starts. As a result, interleaving
of operations take place due
to which T_2 reads the correct
value of x but the incorrect

Value of Y and sum computed by

$$T_2(C+x) = 50,000 + 500 =$$

$$50,500$$

is thus not consistent with the total sum at end of the transaction.

$$T(C+x) = 5,000 + 450 = 50,450$$

This results in database inconsistency due to a loss of 50 units. Hence transaction must take place in isolation and changes should be visible only after they have been made to the main memory.

Durability :-

This property ensures that once the transaction has completed execution, the updates and modifications to the database

are stored in and written to disk
and may persist even if a
system failure occurs. These
updates now become permanent
and are stored in non-volatile
memory. The effects of the
transaction, thus are never lost.

(92 1021 1000 2 1111 X)

do not update to others until
hosted by a number of sub
units for normal and
blood experiments maintains
the right types of address
and responses of alarm and

alarms. The queue will
be used to maintain the order
of topology, first with nodes
linked with a linked list.

Concurrent Executions

concurrent execution refers to the simultaneous execution of more than one transaction.

There are two good reasons for allowing concurrency

- ① Improved throughput and resource utilization! - A transaction may contain I/O activity or CPU activity. The CPU and I/O activity in computer system can operate in parallel.

Read or write operations of one transaction is in progress on one disk, another transaction can be running in the CPU, while another disk may be executing a read or write on behalf of a third transaction. All of this increases the

throughput of the system. Correspondingly, the processor and disk utilization also increases.

(b) Reduced waiting time !

If the transactions are operating on different parts of the database it is better to let them run concurrently, sharing the CPU cycles and disk access among them.

Here schedule 1 - a serial schedule in which T_1 is followed by T_2

	T_1	T_2
	read(A)	
	$A := A - 50$	
950	write(A)	
	read(B)	
	$B := B + 50$	
2050	write(B)	
		read(A)
		$\text{temp} := A * 0.1$

$$\begin{aligned} A &= 100 \\ B &= 2000 \end{aligned}$$

(2)

T_1	T_2	
	$A := A - \text{temp}$	$A = 855$
	$\text{write}(A)$	$@405$
	$\text{read}(B)$	$2050 + 95$
	$B := B + \text{temp}$	2145
	$\text{write}(B)$	

so final $A = 855$ $B = 2145$

Suppose the current values of account A and B are \$1000 and 2000 respectively. Then Transaction

T_1 transfers \$500 from A to account B. And Transaction T_2 transfers 10% of balance from account A to account B.

The final values of accounts A and B after execution of T_1 and T_2 transactions in order are 855 and 2145 respectively.

Sum $855 + 2145 = 3000$ is preserved after the execution of both transaction.

Similarly if the transactions are executed one after another in order T_2 followed by T_1 sum again expected, the sum $A + B$ is preserved.

"C Schedule - a serial schedule in which T_2 after T_1)

T_1	T_2
read(A)	
	temp = A * 0.1
	$A := A - \text{temp}$
	write(A)
	Read(B)
	$B := B + \text{temp}$
	write(B)
read(A)	
$A := A - 50$	
write(A)	
read(B)	
$B := B + 50$	
	write(B)

(8)

Schedule 3 - a concurrent schedule
equivalent to Schedule 1

	T_1	T_2
$A = 1000$		
$B = 2000$		
	read(A)	
	$A := A - 50$	
	write(A)	
$A = 950$		
		read(A)
		$\text{temp} := A * 0.1$
		$A := A - \text{temp}$
		write(A) $A = 855$
	read(B)	
$B = 2050$	$B := B + 50$	
	write(B)	
		read(B)
		$B := B + \text{temp}$
		write(B) $B = 2145$

Schedule 3 - a concurrent schedule
equivalent to Schedule 1.

$$855 + 2145 = 3000$$

$$A = 100^{\circ}$$

$$B = 200^{\circ}$$

$$A = 950$$

$$A = A - 50$$

T_2

$$\text{read}(A)$$

$$\text{temp} = A + 0.1 \quad A = 95$$

$$A = A - \text{temp}$$

$$\text{write}(A) \quad A = 855$$

$$\text{read}(B) \quad B = 2000$$

$$A = 950$$

$$\text{write}(A)$$

$$2000$$

$$\text{read}(B)$$

$$2000 + 50$$

$$= 2050$$

$$B = B + 50$$

$$\text{write}(B)$$

$$B = B + \text{temp} \quad 2050 + 95$$

$$\text{write}(B) \quad B = 2145$$

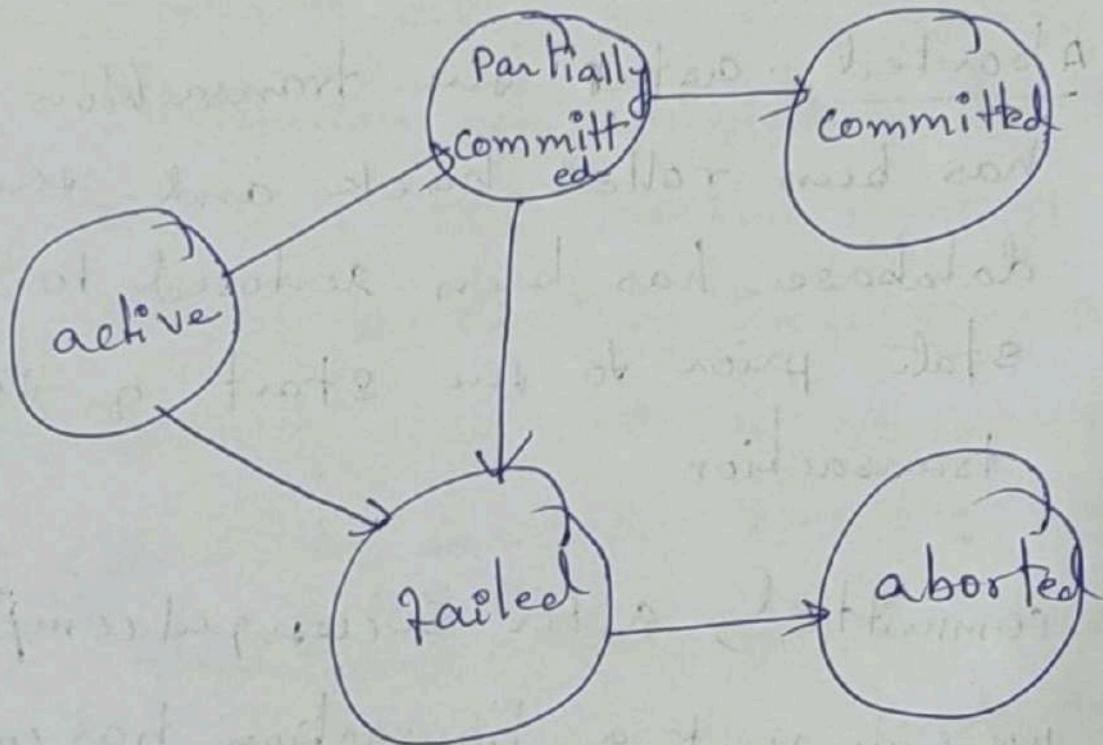
Schedule 4

After the execution of this schedule, we arrive at a state where the final accounts A and B are 950 and 2145 respectively. So final state is an inconsistent state.

$$950 + 2145 = 3095 \text{ inconsistent state.}$$

Transaction state

state diagram of a transaction.



A transaction must be in one of the following states:

- ① Active! - the initial state;
the transaction stays in this state while it is executing
- ② Partially committed; after the final statement has been executed

③ Failed :- after the discovery that normal execution can no longer proceed.

④ Aborted, after the transaction has been rolled back and the database has been restored to its state prior to the start of the transaction.

⑤ committed, after successful completion we say that a transaction has committed only if it has entered the committed state. Similarly, we say that a transaction has aborted only if it has entered the aborted state. A transaction is said to have terminated if it has either committed or aborted.

A transaction starts in the active state. When it finishes its final statement, it enters the partially committed state. At this point, the transaction has completed its execution, but it is still possible that it may have to be aborted, since the actual output may still be temporarily residing in main memory and thus if a hardware failure occurs, the transaction ~~as~~ ^{leads to} may not be successfully complete.

The database system then writes out enough information to disk such that, even in the event of a failure, the updates performed by the transaction can be re-created when the system restarts after the failure. When the last of this

information is written out, the transaction enters the committed state

Implementation of Atomicity and Durability

- The recovery - management component of a database system can support atomicity and durability by a variety of schemes.

shadow

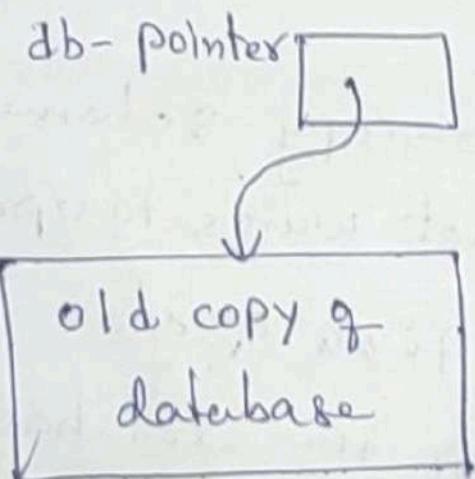
shadow copy scheme is the simplest scheme.

- In the shadow-copy scheme, a transaction that wants to update the database first creates a complete copy of the database. All updates are done on the new database copy, leaving the original copy, the shadow copy, untouched. If at any

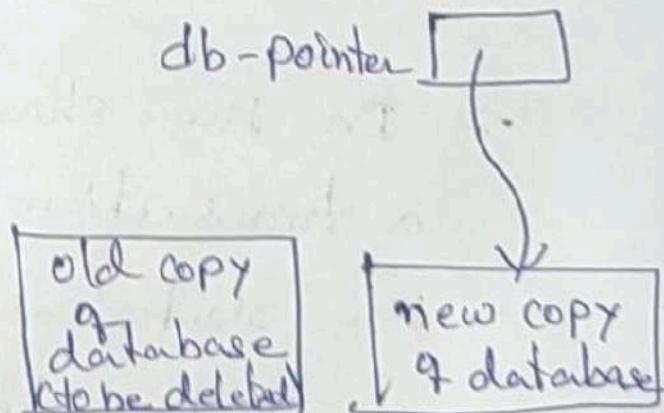
Point the transaction has to be aborted, the system deletes the new copy. The old copy of the database has not been affected.

The scheme also assumes that the database is simply a file on disk.

A pointer called db-pointer is maintained on disk. It points to the current copy of the database.



(a) Before update



(b) After update

shadow-copy technique for atomicity and durability.

(2)

If the transaction completes, it is committed as follows.

- (a) First, the operating system is asked to make sure that all pages of the new copy of the database have been written out to disk.
- (b) After the operating system has written all the pages to disk, the database system updates the pointer db-pointer to point to the new copy of the database. The new copy then becomes the current copy of the database. The old copy of the database is then deleted.

Implementation & Isolation

A DBM's isolation feature ensures that several transactions can take place simultaneously and that no data from one database should have an impact on another. In other words, the process on the second state of the database will start after the operation on the first state is finished.

Implementation of isolation typically involves concurrency control mechanisms. Here are common mechanisms used

①

Locking Mechanisms

Locking ensures exclusive access to a data item for a transaction. This means that while one transaction holds a lock on a data item, no other transaction can access that item.

- Shared Lock (S-Lock): Allows a transaction to read an item but not write to it.
- Exclusive Lock (X-Lock): Allows a transaction to read and write an item. No other transaction can read or write until the lock is released.

→ Two phase Locking (2PL).

This protocol ensures that a transaction acquires all the locks before it releases any. This results in a growing phase (acquiring locks and not releasing any) and a shrinking phase (releasing locks and not acquiring any).

2- Timestamp-based protocols

Every transaction is assigned a unique timestamp when it starts.

This timestamp determines the order of transactions. Transactions can only access the database if they respect the timestamp order, ensuring older transactions get priority.

Serializability

The database serializability is a concept that helps to identify which non-serial schedules are correct and will maintain the consistency of the database. Because all non-serial schedules are not correct, some may lead to inconsistency of data base.

There are two types of serializability

- ① conflict serializability
- ② view serializability

Conflict Serializability :-

In a given non-serial schedule can be converted into a serial schedule by swapping its

non conflicting operations then it is called as conflict serializable.

schedules.

→ while swapping the order of execution of two conflicting operations cannot be changed because if they are applied in different order, they can have different effect on the database or on the other transactions in the schedule.

T ₁	T ₂
read(A) write(A)	read(A) write(A)
read(CB) write(CB)	read(CB) write(B)

Conflict Serializability

Let us consider a schedule s in which there are two consecutive instructions I_i and I_j of transactions T_i and T_j respectively ($i \neq j$).

- ① If I_i and I_j refer to different data items, then we can swap I_i and I_j without affecting the results of any instruction in the schedule.
- ② But if I_i and I_j refer to the same data item o , then the order of the two steps may matter.
Since we are dealing with only read and write instructions there are four cases that we need to consider.
 - (a) $I_i = \text{read}(o)$ and $I_j = \text{read}(o)$

The order of I_i and I_j does not

matter, since the same value of a is read by T_i and T_j regardless of the order.

(b) $I_j = \text{read}(a)$, $I_i = \text{write}(a)$. If I_i comes before I_j , then T_i does not read the value of a that is written by T_j in Instruction I_j . If I_j comes before I_i , then T_i reads the value of a that is written by T_j . Thus the order of I_i and I_j matters.

(c) $I_i = \text{write}(a)$, $I_j = \text{read}(a)$
The order of I_i and I_j matters for reasons similar to those of the previous case.

a) $D_i = \text{write}(o), I_j = \text{write}(o)$

Since both instructions are write operations, the order of these instructions does not effect either T_i or I_j .

If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are conflict equivalent.

The concept of conflict equivalence leads to the concept of conflict serializability we say that a other schedule S is conflict serializable if it is conflict equivalent to a serial schedule.

S1	
T1	T2
read(A)	
write(A)	
	read(A)
	write(A)
read(CB)	
write(CB)	

Swapping

T1	T2
read(CB)	(Schedule 1)
write(B)	

Here write(A) of T2 and read(CB) of T1

we can swap as there is no

conflict because of different data

type	
T1	T2
read(A)	
write(A)	
	read(A)
red(CB)	
	write(CA)
write(CB)	

swapping

(Schedule 2)

T1	T2
read(A)	
write(A)	
read(B)	
	read(A)
	write(A)
write(B)	

swapping

read(CB)
write B

so schedule ① and schedule 2
both produce the same final
system state.

So we say that Schedule 1 and Schedule 2 are conflict equivalent

S2

T ₁	T ₂
read CA)	
write (A)	
read (B)	
write (B)	read (A)
	write (A)
	read (B)
	write (B) schedule - 3

Now this Schedule 3 is conflict serializable.

View Serializability

Consider two schedules S and S' where the same set of transactions participates in both schedules.

The schedules S and S' are said to be view equivalent if

three conditions are met

(1) For each data item α , if transaction T_i reads the initial value of α in schedule s , then transaction T_i must, in schedule s' also read the initial value of α .

(2) For each data item α , if transaction T_i executes read(α) in schedule s , and if that was produced by a write(α) operation executed by transaction T_j , then the read(α) operation of transaction T_i must, in schedule s' , also read the value of α that was produced by the same write(α) operation of transaction T_j .

For each data item a , the transaction (T_a) that performs the final write(a) operation in schedule s must perform the final write(a) operation in schedule's.

Condition 1 and 2 ensure that each transaction reads the same values in both schedules and therefore performs the same computation.

Condition 3, coupled with conditions 1 and 2 ensures that both schedules result in the same final system state.

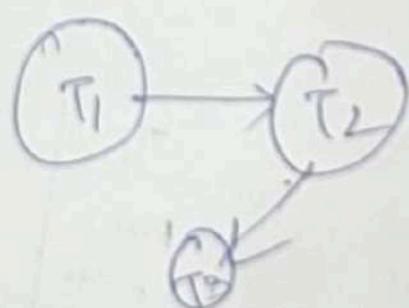
say
 s

$A = 100$			s		
T_1	T_2	T_3	T_1	T_2	T_3
$R(A)$			$R(A)$		
	$A = A - 40$		$A = A - 40$		
	$w(A)$		$w(A)$		
$A = A - 40$			$A = A - 40$		
$w(A)$			$w(A)$		
		$A = A - 20$			
		$w(A)$			
		$A = 0$			
					$A = 0$

Schedule 1		Schedule 2		
T3	T4	T3	T4	T6
read(o)	write(o)	read(o)	write(o)	
write(o)		write(o)		write(o)
				write(o)

Here Schedule 2 is view serializable

so we say schedule s' is view serializable



No Loop

s is view equivalent to s'
and s' is view serializable.

Testing for Serializability

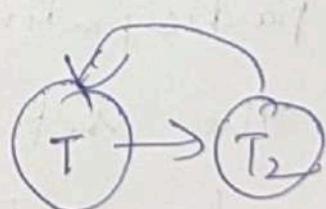
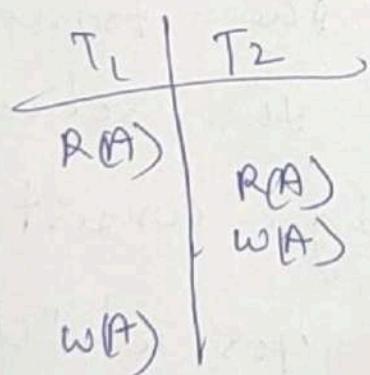
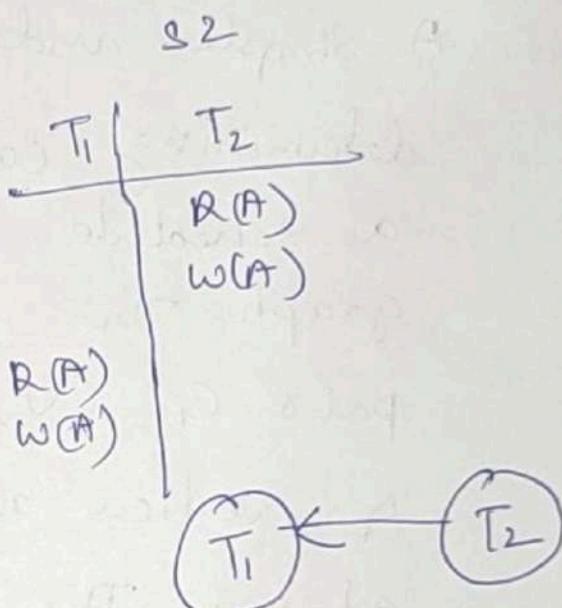
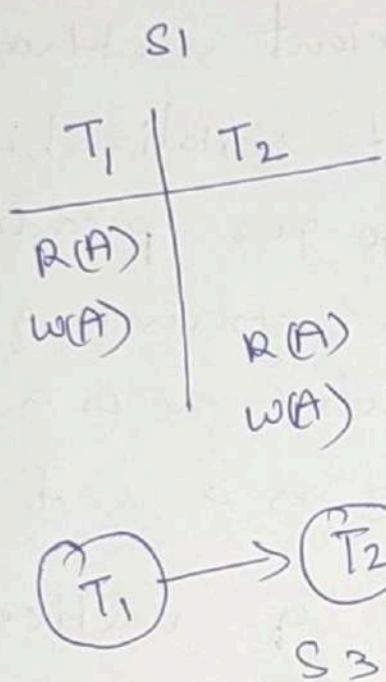
when designing concurrency control, we must show that schedules generated by the scheme are serializable.

A simple and efficient method for determining conflict serializability of a schedule. we go for procedure graph. The graph consists of a pair $G = (V, E)$ where V is a set of vertices and E is a set of edges. The set of vertices consists of all the transactions participating in the schedule. The set of edges consist of all edges $T_i \rightarrow T_j$ for which one of three conditions holds:

- ① T_i executes $\text{write}(o)$ before T_j executes $\text{read}(o)$.

2. T_i executes $\text{read}(c)$ before T_j executes $\text{write}(c)$.

③ T_i executes $\text{write}(c)$ before T_j executes $\text{write}(c)$.



If the precedence graph for S_3 has a cycle then schedule S is not conflict serializable.

(2)

If the graph contains no cycle,
then the schedule is conflict
serializable.

Concurrency control

Concurrency control is a crucial database management system (DBMS) component.

It manages simultaneous operations without them conflicting with each other. The primary aim is maintaining consistency, integrity and isolation when multiple users or applications access the database simultaneously.

Problems with concurrent execution

In a database transaction, the two main operations are read and write operations. so there is a need to manage these two operations in the concurrent execution.

Problem ① Lost update problem (W-W conflict)

The problem occurs when two different database transactions perform one read/write operations on the same database items in an interleaved manner (in concurrent execution) that makes the values of the items incorrect hence making the database inconsistent.

For example

Consider the below diagram where two transactions T_x and T_y are performed on the same account A where the balance of account A is 300

Time	T_x	T_y	let $A = 300$
t_1	Read(A)	-	
t_2	$A = A - 50$	Read(A)	
t_3	-	$A = A + 100$	
t_4	-	-	
t_5	-	-	
t_6	Write(A)	-	
t_7	-	Write(A)	

- At time t_1 , transaction T_x reads the value of account A is ₹ 300
- At time t_2 , transaction T_x deducts 50₹ from account A that becomes ₹ 250
- Alternatively at time t_6 , transaction T_y reads the value of account A that will be ₹ 300

only because Tx didn't update the value yet.

- At time t_4 transaction Tx adds Rs 100 to account A that becomes 400.
- At time t_6 , transaction Tx writes the values of account A that will be updated as 250 as only as Tx didn't update yet.
- Similarly at time t_7 , transaction Tx writes the values of account A, so it will write as done at time t_4 that will be \$400.
It means the value written by Tx is lost as Rs 250 is lost.
Hence data becomes incorrect, and database sets to inconsistent.

②

3 Dirty Read problems (W-R conflict)

The dirty read problem occurs when one transaction updates an item in the database, and somehow the transaction fails, and before the data gets rolled back, the updated database item is accessed by another transaction.

Then comes write-read conflict between both transactions.

Poor example.

Consider two transactions T_x and T_y in the below diagram performing read/write operations on account A where the available balance is account A is 300

Time	T_x	T_y
t_1	Read(A)	-
t_2	$A = A + 50$	
t_3	write(A)	
t_4	-	Read(A)
t_5	Server down Roll back	-

- At time t_1 , transaction T_x reads the value of account A, ie 300
- At time t_2 , transaction T_x adds 50 to account A that becomes 350
- At time t_3 , transaction T_x writes the updated value in account A ie "350" $\underline{\text{Rs}}$
- Then at time t_4 , transaction T_y reads account A that will be read as 350 Rs

- Then at time t_5 , transaction T_2 rolls back due to server problem, and the value changes back to 300 (as initially).
- But the value for account A remains 350 as for transaction T_2 as committed, which is the dirty read and therefore known as Dirty Read problem.

(3)

Unrepeatable Read problem.

(W-W conflict)

Also known as inconsistent retrieval problem that occurs when in a transaction, two different values are read for the same database item.

Consider two transaction T_2 and T_3 performing the read/write

operations on account A, having
an available balance = 300 Rs

Time	Tx	Ty
t_1	Read(A)	-
t_2	-	Read(A)
t_3	-	$A_2 = A + 100$
t_4	-	write(A)
t_5	Read(A)	-

- At the time t_1 , transaction Tx reads the value from account A i.e 300 Rs
- At the time t_2 , transaction Ty reads the value of account A, it is 300.
- At the time t_3 , Transaction Ty updates the value of account A by adding 100 Rs to the available

balance and then it becomes 400

- At time t_4 transaction T_y writes the updated value in Rs 400
- After that, at time t_5 , transaction T_x reads the available value of account A, and that will be read as 400.
- It means that within the same transaction T_x , it reads two different values of account A, in 300 initially, and after updation made by transaction T_y , it reads 400Rs. This is an unrepeatable read and is therefore known as the unrepeatable read problem.